


Name: Divyank Kulshrestha

Cluster Colors

The output shows text of different colors as white on a black background. I sub sampled the image in order to reduce computations. And then displayed each k-means cluster as white on a black background image. I used 10 clusters, instead of 8, to account for all the colors in the image, because using 8 clusters was not able to distinguish between green and blue inks. Some of the results had noise in them. I tried to clean up the noise using morphological operations, but it also erased the desired text, so I had to abort. I also did not try seeding, as the output I got was good enough without it.

OUTPUT IMAGES:



THIS IS A
SECRET MESS-
AGE. IT IS
HIDDEN BEHIND
OTHER COLORS.

RED IS A DIST-
RACTION COLOR.

RED IS QUICKLY
READ IF YOU READ
RED. BE WELL

READ. READ TO
LEAD.

THANK YOU FOR ^{READING} US.

BLUE INK
BLUE
BLUE

BLUE BLUE

BLUE
BLUE

BLUE

BLUE
BLUE
BLUE

BLUE.

GREEN

GREEN

GREEN

GREEN

GREEN.

GREEN


GREEN

GREEN

Find Raspberries

I sub-sampled the image by skipping every fourth row/column for faster computations before getting clicks from the user using `ginput()`. Then I computed the mahalanobis distance for the image using the function '**mahal**'. To find a threshold value, I started with 5 which was producing insufficient results. Finally I tried 10 as the threshold, which seemed to give the most desired output. As I was still a little curious, I also tried using 20 for the threshold, but that ended up detecting the majority of the image as raspberries, which was again incorrect. I displayed the output as a binary image, where white color represents the pixels within the threshold distance (the raspberries in the image).


OUTPUT IMAGE:



Find Oranges

First I rotated the image to align it straight and then sub-sampled the image by skipping every tenth row/column for faster computations. The rest of the procedure is the same as for finding raspberries, with the threshold value used as 10 again.

OUTPUT IMAGE:



Conclusion

I realized the importance of converting the image to double and sub-sampling in this homework. Using K-means takes computations, and sub-sampling helps perform K-means much faster. Using K-means without sub-sampling took a lot of time to finish. Using the function **ginput()** again served as a refresher on one of the ways to get input from the user.

It took me some time to figure out how to **mahal()** function with the user-clicked inputs. But once I was able to calculate that, I used a zeros vector as the black image (pixel value 0), and updated pixels in that to white (pixel value 1), based on whether a pixel is within the threshold distance of 10 from the user-clicked input pixels. I used the same procedure to detect both raspberries as well as oranges.